

vLLM V1: one mixed prefill-decode batch becomes paged FlashAttention inputs

Pinned commit 4c2d2776 Example: causal self-attention, one KV-cache group, no padding/cascade, DCP = 1, non-hybrid pages, $B = 16$

